

Transformer test loss vs. learning rate and depth, at  $n = 1024$

